

Dávid Szakállas

Sr. Software Engineer



Budapest, Hungary



+36 30 626 6558



<https://github.com/dszakallas>



david.szakallas@gmail.com

About

Dávid is an experienced software engineer who focuses on developing data-heavy applications on public clouds and Kubernetes. He is keen on delivering automated, well-engineered solutions to a variety of problems. He is a programming language polyglot and tries to match the tools to the job. In his past roles he lead the migration of workloads from bare VMs in the cloud to greenfield Kubernetes clusters and lead architecting and productionizing the Data Lake at Turbine, worked in the team operating the data processing platform at Hiya, developed a supervised ML pipeline for credit fraud detection and an identity graph database at Ekata. He holds an M.Sc. in Computer Science Engineering.

Skills

Scala



Apache Spark



Python



Kubernetes Application Development



Data Lake



AWS



DevOps



Kubernetes Administration



Azure



Functional Programming



Linux Containers



Experience

Now – June 2021 Staff Software Engineer

Turbine [cancer drug discovery | early-stage startup]

Developing our data lake on Azure with Databricks, Airflow and ADLSv2. Productionizing existing ETL workflows. Leading Spark adoption, establishing DevOps practices for data engineers. Advocating containerization and IaC. Leading the migration of the cell simulation orchestration platform based on Airflow to a greenfield AKS cluster. Establishing GitOps using Flux.

Now – Oct 2020 Individual Open Source Contributor (Part-time contract)

Linked Data Benchmark Council

I am contributing to LDDB's open source Social Network Benchmark data generating framework.

May 2021 – June 2020 Senior Data Engineer

Hiya [mobile spam blocking | mid-stage startup]

Developing and operating our data lake based on AWS Glue, Apache Spark, Apache Airflow, Kubernetes. Automating internal Spark releases tailored to AWS Glue, Kubernetes and in-house libraries. Migrating to Spark 3.0. Developing metrics collection for Spark running in Kubernetes with Prometheus. Developing in-house libraries in Scala and Python.

Feb 2020 – Oct 2017 Data Engineer

Ekata [identity verification | late-stage startup]

Developing the multi-terabyte batch data ingestion pipeline for our proprietary social graph database in Apache Spark on Databricks with complex graph-specific application logic written in Scala / SparkSQL. Gaining experience in machine learning and XGBoost. Orchestrating supervised ML pipelines with Apache Airflow, Kubernetes and Databricks. Developing an ML feature prototyping application and library for data scientists.

May 2017 – Jan 2016 Software Engineer

RisingStack [application monitoring | early-stage startup]

Developing the software application monitoring agent in Node.js. Developing the back-end services for the distributed tracing feature. Gaining experience in distributed tracing technologies such as OpenZipkin and OpenTracing.

Education

2018 – 2015 M.Sc. in CS Eng. Budapest University of Technology and Economics
Thesis: Evaluation of openCypher graph queries with a local search-based algorithm (written in Clojure)

2015 – 2011 B.Sc. in CS Eng. Budapest University of Technology and Economics

Other information

Hobbies

I sometimes program as a hobby in Rust or Clojure. My main editor is Emacs, and my developer notebook OS is either macOS or Arch Linux. I have a programming blog on szakallas.eu. In my past time, I enjoy swimming and scuba diving.

Talks

Oct 2018 Spark+AI Summit EU London – Spark Schema for Free

Languages

Hungarian	Native language
English	Professional working proficiency
German	Basic proficiency