

Dávid Szakállas

SRE & Platform

 Prague, Czechia

 +420 731 886 656

 david.szakallas@gmail.com

 dszakallas

 dszakallas

 szakallas.net

About

Dávid is an experienced engineer specializing in developing and operating highly available systems, cloud platforms, and data pipelines on public clouds and Kubernetes. With a strong background in site reliability and distributed systems, he delivers robust infrastructure, observability, and automation. He holds an M.Sc. in Computer Science Engineering.

Skills

Software Engineering

Python	5.5 y
Scala	5.3 y
Go	3.8 y
Clojure	3.0 y

Site Reliability

CloudOps & Reliability	4.4 y
Observability (Prometheus, ELK)	3.3 y
Incident Response	2.4 y
CloudSec & NetSec (SPIFFE)	2.4 y

DevOps & Platform

Kubernetes (Apps & Platform)	7.3 y
Linux Containers	7.3 y
GitOps (Flux, Argo)	5.8 y
IaC & Terraform	4.4 y
Nix/NixOS	2.0 y

Public Clouds

AWS	5.3 y
Azure	4.9 y

Data Engineering

Apache Spark	6.4 y
Apache Airflow	4.4 y
AWS EMR & Data Lakes	4.0 y
Databricks	3.5 y

Misc

Citizenship
EU (Hungarian)

Languages
Hungarian (Native), English (Professional), Czech (Basic), German (Basic)

EXPERIENCE

Everpure (formerly Pure Storage)

Apr 2024 – Present

Staff Site Reliability Engineer

Developing and operating Everpure Cloud Azure Native. Defining SLOs/SLAs, region provisioning, building observability and alerting pipelines, incident response, and internal developer tooling.

azurek8sgogitopsiacpythonterraformELK stackSPIFFEvclustergithubircloudseccloudopsdevopsnetsecobservability

Everpure (formerly Pure Storage)

Nov 2022 – Apr 2024

Staff Software Engineer

Developed backend services and workflows for SaaS/PaaS control planes written in Go and running on Kubernetes.

gok8sazuregitopsbackend engineeringkubernetes operator development

Turbine

May 2021 – June 2022

Staff Software / Data Engineer

Evaluated and led adoption of Kubernetes (AKS) and GitOps (Flux). Productionized Apache Spark ETL workloads on Databricks with CI/CD and unit tests, and performance tuned PySpark applications.

k8sgitopspythonazurecontainersairflowdatabricksdata lakesdata engineeringplatform engineering

Graph Data Council

Oct 2020 – Sep 2022

Open Source Developer (Part-time contract)

Contributed to LDBC's Social Network Benchmark data generator in Scala and Apache Spark, optimizing generation of multi-terabyte datasets on EMR clusters.

awsemrapache sparkapache hadoopscaladatabase benchmarking

Hiya

June 2020 – Apr 2021

Senior Data Engineer

Developed and operated data lake on AWS Glue, Spark, Airflow, and Kubernetes. Automated Spark releases, built Prometheus metrics collection, and developed in-house libraries in Scala and Python.

scalaapache sparkawsgitopsprometheussoftware engineeringdata lake designdata platforms opsobservability

Ekata

Oct 2017 – Feb 2020

Data Engineer

Developed the batch data ingestion pipeline for proprietary social graph database in Apache Spark, cutting ingestion time by 80%. Orchestrated supervised ML pipelines with Airflow, Kubernetes, and Databricks.

scalak8sairflowxgboostawsdatabrickssoftware engineeringgraph databasesML training and inferenceanomaly detection

EDUCATION

Budapest Univ. of Technology & Economics

2015 – 2018

M.Sc. in Computer Software Engineering

Thesis: openCypher graph query evaluation with a local search-based algorithm (in Clojure)

Budapest Univ. of Technology & Economics

2011 – 2015

B.Sc. in Computer Software Engineering

PROJECTS

Hobby & Open Source

-  mixxx-launchpad
-  bikeshed
-  podcaster
- Home Server
- Retrogaming

HID mappings for open-source DJ software (Mixxx) 

Nix utilities for workstation configuration 

Podcast generation with Google NotebookLM  

Self-hosted media cluster on NixOS & k3s    

NixOS modules for emulation on Raspberry Pis   

Publications & Talks

2023
Oct 2018

LDBC: Driving competition and collaboration in graph data
Spark+AI Summit EU London Spark Schema for Free